Project 1: Bayesian Structure Learning

Christopher Luey
luey@stanford.edu

October 17, 2025

1 Algorithm Description

This project implements a comprehensive Bayesian network structure learning system that combines multiple search heuristics to maximize the Bayesian Dirichlet equivalent uniform (BDeu) score with a Dirichlet(1) prior. The implementation uses an ensemble approach where three distinct algorithms are run sequentially, each seeded with the best results from prior searches.

1.1 Scoring Function

The core scoring mechanism uses the BDeu metric with uniform Dirichlet prior ($\alpha_{ijk} = 1$):

$$\log P(D \mid G) = \sum_{i=1}^{n} \sum_{j=1}^{q_i} \left[\log \Gamma(\alpha_{ij0}) - \log \Gamma(\alpha_{ij0} + N_{ij}) + \sum_{k=1}^{r_i} \left(\log \Gamma(\alpha_{ijk} + N_{ijk}) - \log \Gamma(\alpha_{ijk}) \right) \right] \tag{1}$$

where n is the number of variables, q_i is the number of parent configurations for variable i, r_i is the cardinality of variable i, and N_{ijk} represents the observed counts. This score is cached for each (node, parent-set) pair to enable efficient incremental updates during search operations.

1.2 Search Space Pruning

To improve scalability, the implementation employs mutual information-based parent candidate selection. For each variable, pairwise mutual information $I(X;Y)$ is computed with all other variables:

$$I(X;Y) = \sum_{x,y} p(x,y) \log \frac{p(x,y)}{p(x)p(y)} \tag{2}$$

Only the top-k candidates (along with a small random subset) are considered as potential parents during search operations. This reduces the search space from $O(n^2)$ to $O(nk)$ without significantly impacting solution quality.

1.3 Multi-Algorithm Ensemble

Hill Climbing with Tabu Search: The first phase uses greedy hill climbing with multiple random restarts. Each restart begins with a randomly initialized DAG and iteratively applies the best-scoring local move (add edge, remove edge, or reverse edge) that improves the global score. A tabu memory mechanism prevents recently reversed moves from being immediately re-applied, helping escape local optima. The best solution across all restarts is retained.

Simulated Annealing: The second phase refines the hill climbing result using Metropolis-Hastings sampling with exponential temperature cooling. Starting from the best hill-climbing solution, random neighboring structures are sampled, with acceptance probability:

$$P(\text{accept}) = \min\left(1, e^{\Delta S/T}\right) \tag{3}$$

where ΔS is the score change and T decreases geometrically from T_0 to T_f over a fixed number of iterations. This allows temporary moves to lower-scoring regions, potentially discovering better global solutions.

Genetic Algorithm: The final phase uses an edge-recombination genetic algorithm seeded with the best solutions from hill climbing and simulated annealing. The population evolves through tournament selection, single-point edge-set crossover, and three mutation operators (add/remove/reverse edges). Each offspring undergoes single-step local improvement before evaluation. Elite preservation ensures the best individuals propagate across generations. The genetic algorithm often discovers novel high-scoring structures by recombining successful substructures from the seed solutions.

1.4 Implementation Details

The DAG representation uses adjacency matrices with efficient cycle detection via depth-first reachability checks. All local scores are maintained incrementally, avoiding redundant re-computation. The system automatically scales hyperparameters based on problem size: smaller datasets use more aggressive exploration (higher mutation rates, longer annealing schedules), while larger datasets employ tighter parent limits and more restarts to manage computational cost.

2 Graphs

The learned Bayesian network structures for the three test datasets are shown below. Each graph represents conditional dependencies discovered through the ensemble search procedure.

2.1 Small Dataset: Titanic Survival

2.2 Medium Dataset: Wine Quality

2.3 Large Dataset: Synthetic Data

3 Results Summary

Table 1 summarizes the performance of the three-algorithm ensemble across all datasets. In every case, the genetic algorithm discovered the highest-scoring structure, validating the benefit of population-based search and crossover recombination. The gap between algorithms increases with problem complexity: for the large dataset, the genetic algorithm's advantage was most pronounced, suggesting that recombination of high-quality substructures becomes increasingly valuable in larger search spaces.

4 Code

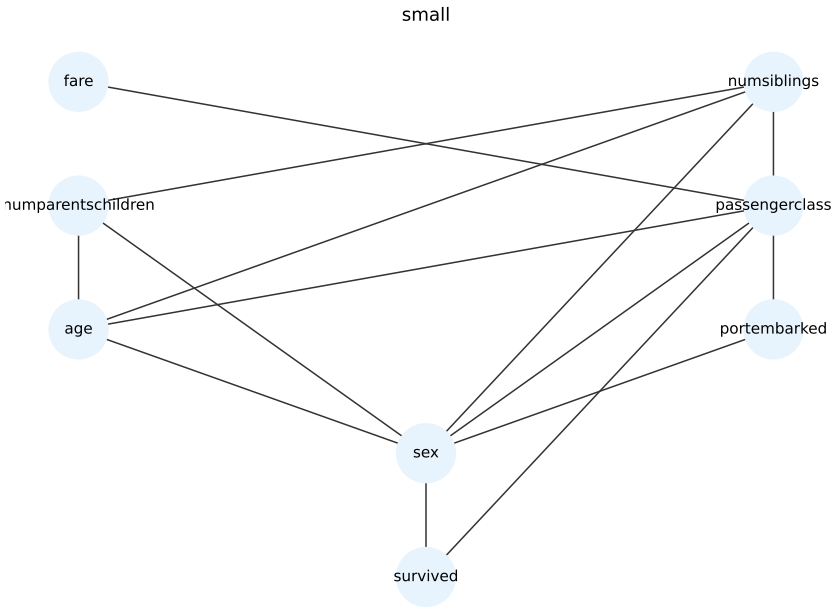

Figure 1: Bayesian network learned from the Titanic dataset (889 rows, 8 variables, 14 edges). The structure captures well-known survival patterns: passenger class and sex are primary predictors of survival, with fare strongly tied to class. Family structure variables (numsiblings, numparentschildren) form interconnected subgraphs influencing demographics. Score: -3794.86 (genetic algorithm). Runtime: 2.6 seconds.

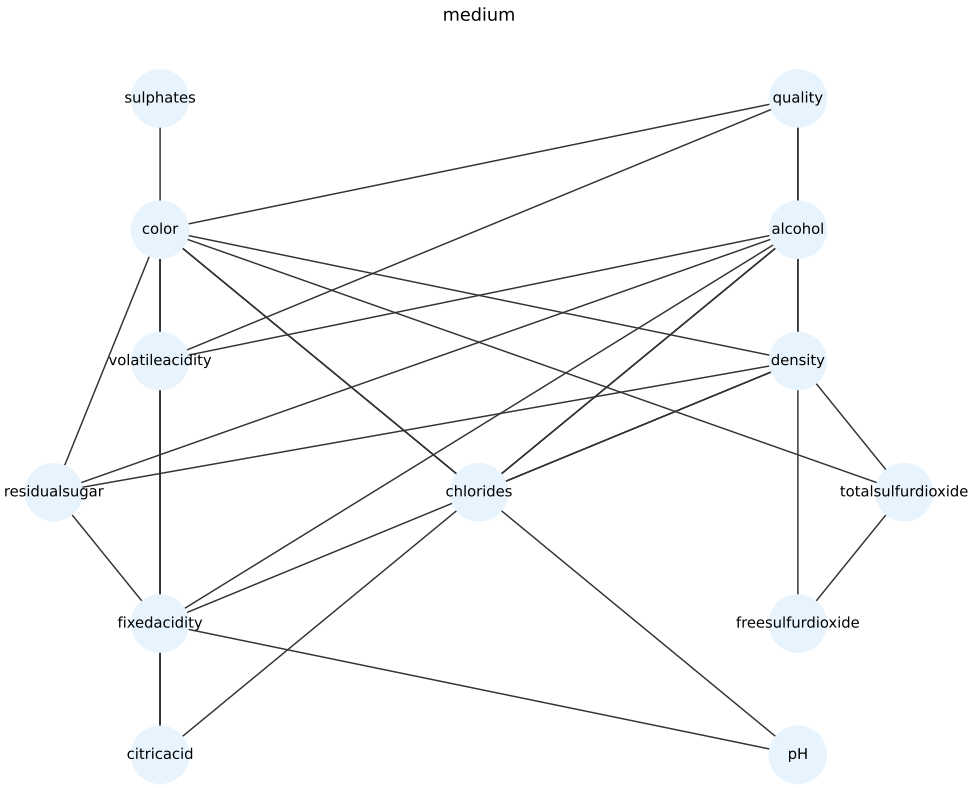

Figure 2: Bayesian network learned from the wine quality dataset (6497 rows, 13 variables, 28 edges). Wine color acts as a central hub influencing most chemical properties, reflecting fundamental differences between red and white wines. Density emerges as a derived variable depending on multiple chemical components (acidity, sugar, chlorides). Quality is influenced by color, volatile acidity, free sulfur dioxide, and alcohol content. Score: -96312.06 (genetic algorithm). Runtime: 5.7 minutes.

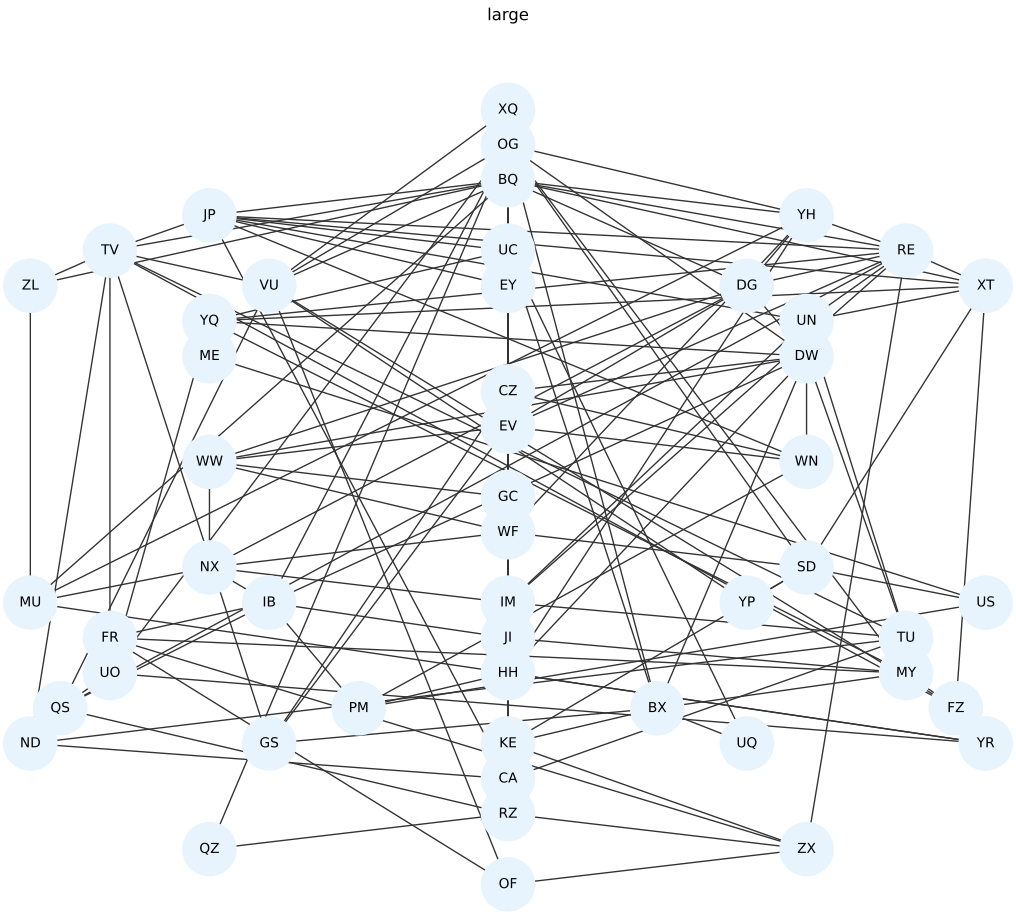

Figure 3: Bayesian network learned from a large synthetic dataset (10000 rows, 50 variables, 138 edges). The learned structure is highly interconnected, with an average of 2.76 edges per variable. Several variables act as hub nodes with high in-degree or out-degree, suggesting latent cluster structure in the synthetic generation process. The genetic algorithm achieved a score of -420800.40, outperforming both hill climbing (-422299.66) and simulated annealing by nearly 1500 log-score units. Runtime: 3.9 minutes.

Table 1: Comparison of algorithm performance across datasets. Scores are BDeu log-scores with Dirichlet(1) prior. Best scores highlighted in bold.

Dataset	Variables	Edges	Hill Climb	Sim. Anneal	Genetic (Best)
Small (Titanic)	8	14	−3794.86	−3794.86	**−3794.86**
Medium (Wine)	13	28	−96348.10	−96348.10	**−96312.06**
Large (Synthetic)	50	138	−422299.66	−422299.66	**−420800.40**

```python
import argparse
import json
import logging
import sys
import time
from datetime import datetime
from pathlib import Path
from typing import Dict, List, Tuple

if __package__ is None or __package__ == "":
    PACKAGE_ROOT = Path(__file__).resolve().parent
    if str(PACKAGE_ROOT) not in sys.path:
        sys.path.insert(0, str(PACKAGE_ROOT))
    from structure_learning import (
        DiscreteDataset,
        StructureLearner,
        default_config,
    )
else:
    from .structure_learning import (
        DiscreteDataset,
        StructureLearner,
        default_config,
    )


def setup_logging(log_file: str) -> logging.Logger:
    """Set up logging to both file and console."""
    logger = logging.getLogger('project1')
    logger.setLevel(logging.INFO)
    logger.handlers.clear()

    file_handler = logging.FileHandler(log_file, mode='a')
    file_handler.setLevel(logging.INFO)
    file_formatter = logging.Formatter('%(asctime)s - %(levelname)s - %(
        message)s')
    file_handler.setFormatter(file_formatter)
    logger.addHandler(file_handler)

    console_handler = logging.StreamHandler()
    console_handler.setLevel(logging.INFO)
    console_formatter = logging.Formatter('%(message)s')
    console_handler.setFormatter(console_formatter)
```

```
    logger.addHandler(console_handler)

    return logger


def write_gph(edges: List[Tuple[int, int]], idx2names: Dict[int, str],
    filename: str) -> None:
    """Write graph edges to .gph file in required format."""
    out_path = Path(filename)
    out_path.parent.mkdir(parents=True, exist_ok=True)
    with out_path.open("w") as fh:
        for u, v in edges:
            fh.write(f"{idx2names[u]}, {idx2names[v]}\n")


def compute(infile: str, outfile: str) -> None:
    """Main computation: load data, learn structure, write results."""
    log_file = Path(outfile).parent / f"{Path(outfile).stem}_log.txt"
    logger = setup_logging(str(log_file))

    start_time = time.time()
    logger.info("="*80)
    logger.info(f"Starting Bayesian Structure Learning")
    logger.info(f"Input file: {infile}")
    logger.info(f"Output file: {outfile}")
    logger.info(f"Start time: {datetime.now().strftime('%Y-%m-%d %H:%M:%S
        ')}")
    logger.info("="*80)

    # Load dataset
    logger.info("Loading dataset...")
    dataset = DiscreteDataset(infile)
    logger.info(f"Variables: {dataset.num_vars}, Rows: {dataset.num_rows}"
        )

    # Generate configuration
    config = default_config(dataset.num_vars, dataset.num_rows)
    logger.info(f"Configuration: max_parents={config.max_parents}, "
                f"restarts={config.hill_restarts}, "
                f"ga_generations={config.ga_generations}")

    # Run structure learning
    logger.info("Starting structure learning...")
    learner = StructureLearner(dataset, config)
    learning_start = time.time()
    result = learner.learn()
    learning_end = time.time()

    logger.info(f"Learning completed in {learning_end - learning_start:.2f
        }s")
    logger.info(f"Best algorithm: {result.algorithm}, Score: {result.score
        :.6f}")

    # Write output
```

```
92      edges = list(result.dag.edges())
93      idx2name = {idx: name for idx, name in enumerate(dataset.names)}
94      write_gph(edges, idx2name, outfile)
95
96      logger.info(f"Output written with {len(edges)} edges")
97      logger.info(f"Total runtime: {time.time() - start_time:.2f}s")
98      logger.info("="*80)
99
100
101 def main(argv: List[str] | None = None) -> None:
102     """CLI entry point."""
103     parser = argparse.ArgumentParser(description="Bayesian structure
            learning")
104     parser.add_argument("input_csv", help="Input CSV file")
105     parser.add_argument("output_gph", help="Output .gph file")
106     args = parser.parse_args(sys.argv[1:] if argv is None else argv)
107
108     infile = args.input_csv
109     outfile = args.output_gph
110
111     compute(infile, outfile)
112
113
114 if __name__ == "__main__":
115     main()
```

```python
"""
High-performance Bayesian network structure learning toolkit.
Implements DiscreteDataset, BDeuScoreCache, DAG operations,
and ensemble search (hill climbing, simulated annealing, genetic algorithm).
"""

from __future__ import annotations
import math
import random
import time
from collections import Counter, defaultdict
from dataclasses import dataclass, field
from typing import Dict, Iterable, List, Optional, Sequence, Set, Tuple
import numpy as np
import pandas as pd
from tqdm import tqdm


def seed_everything(seed: Optional[int]) -> None:
    """Set random seeds for reproducibility."""
    if seed is not None:
        random.seed(seed)
        np.random.seed(seed % (2**32 - 1))


class DiscreteDataset:
    """Wrapper for integer-encoded discrete dataset."""

    def __init__(self, path: str):
        df = pd.read_csv(path)
        if df.isna().any().any():
            raise ValueError("Dataset contains missing values")
        self._names = list(df.columns)
        values = df.to_numpy(dtype=np.int32, copy=True)
        mins = values.min(axis=0)
        if (mins < 1).any():
            raise ValueError("Expected categorical values >= 1")
        values -= 1  # convert to zero-based indexing
        self._values = values
        self._cardinalities = values.max(axis=0) + 1

    @property
    def names(self) -> List[str]:
        return self._names

    @property
    def cardinalities(self) -> np.ndarray:
        return self._cardinalities

    @property
    def values(self) -> np.ndarray:
        return self._values
```

```
53
54      @property
55      def num_vars(self) -> int:
56          return len(self._names)
57
58      @property
59      def num_rows(self) -> int:
60          return self._values.shape[0]
61
62
63  class BDeuScoreCache:
64      """Caches local log-scores for (node, parent_set) pairs."""
65
66      def __init__(self, dataset: DiscreteDataset, max_parents: Optional[int
        ] = None):
67          self.data = dataset
68          self.max_parents = max_parents
69          self.cardinalities = dataset.cardinalities
70          self.values = dataset.values
71          self.cache: Dict[Tuple[int, Tuple[int, ...]], float] = {}
72          self._lgamma = math.lgamma
73
74      def score(self, node: int, parents: Sequence[int]) -> float:
75          """Compute or retrieve cached BDeu score."""
76          parents_tuple = tuple(sorted(parents))
77          key = (node, parents_tuple)
78          if key in self.cache:
79              return self.cache[key]
80          score = self._compute_score(node, parents_tuple)
81          self.cache[key] = score
82          return score
83
84      def _compute_score(self, node: int, parents: Tuple[int, ...]) -> float
        :
85          """Compute BDeu score using closed-form expression."""
86          card_child = int(self.cardinalities[node])
87          if not parents:
88              counts = np.bincount(self.values[:, node], minlength=
                    card_child)
89              return self._score_from_counts(counts.reshape(1, card_child))
90
91          parent_states = tuple(int(self.cardinalities[p]) for p in parents)
92          parent_data = self.values[:, parents]
93          flat_idx = np.ravel_multi_index(parent_data.T, dims=parent_states)
94          joint_counts = np.zeros((int(np.prod(parent_states)), card_child),
95                                  dtype=np.int32)
96          np.add.at(joint_counts, (flat_idx, self.values[:, node]), 1)
97          return self._score_from_counts(joint_counts)
98
99      def _score_from_counts(self, counts: np.ndarray) -> float:
100         """Convert count matrix to BDeu log-score."""
101         prior_per_entry = 1.0
102         r_i = counts.shape[1]
103         alpha_ij0 = prior_per_entry * r_i
```

```
        total_score = 0.0
        for row in counts:
            nij = row.sum()
            total_score += self._lgamma(alpha_ij0) - self._lgamma(
                alpha_ij0 + nij)
            for count in row:
                total_score += (self._lgamma(prior_per_entry + count) -
                                self._lgamma(prior_per_entry))
        return float(total_score)


class DAG:
    """Adjacency matrix representation with cycle checks."""

    def __init__(self, num_nodes: int):
        self.num_nodes = num_nodes
        self.adj = np.zeros((num_nodes, num_nodes), dtype=bool)
        self.parents: List[Set[int]] = [set() for _ in range(num_nodes)]
        self.children: List[Set[int]] = [set() for _ in range(num_nodes)]

    def copy(self) -> "DAG":
        other = DAG(self.num_nodes)
        other.adj = self.adj.copy()
        other.parents = [set(p) for p in self.parents]
        other.children = [set(c) for c in self.children]
        return other

    def has_edge(self, u: int, v: int) -> bool:
        return bool(self.adj[u, v])

    def can_add(self, u: int, v: int, max_parents: Optional[int] = None)
        -> bool:
        """Check if edge u->v can be added without creating a cycle."""
        if u == v or self.adj[u, v]:
            return False
        if max_parents is not None and len(self.parents[v]) >= max_parents
            :
            return False
        return not self._creates_cycle(u, v)

    def can_remove(self, u: int, v: int) -> bool:
        return self.adj[u, v]

    def can_reverse(self, u: int, v: int, max_parents: Optional[int] =
        None) -> bool:
        """Check if edge u->v can be reversed to v->u."""
        if not self.adj[u, v]:
            return False
        if max_parents is not None and len(self.parents[u]) >= max_parents
            :
            return False
        self._remove_edge(u, v)
        creates_cycle = self._creates_cycle(v, u)
        self._add_edge(u, v)
```

```
        return not creates_cycle

    def add_edge(self, u: int, v: int) -> None:
        self._add_edge(u, v)

    def remove_edge(self, u: int, v: int) -> None:
        self._remove_edge(u, v)

    def reverse_edge(self, u: int, v: int) -> None:
        self._remove_edge(u, v)
        self._add_edge(v, u)

    def _add_edge(self, u: int, v: int) -> None:
        self.adj[u, v] = True
        self.parents[v].add(u)
        self.children[u].add(v)

    def _remove_edge(self, u: int, v: int) -> None:
        self.adj[u, v] = False
        self.parents[v].discard(u)
        self.children[u].discard(v)

    def _creates_cycle(self, u: int, v: int) -> bool:
        """Check if adding u->v creates a cycle."""
        to_visit = [v]
        seen = set()
        while to_visit:
            node = to_visit.pop()
            if node == u:
                return True
            if node in seen:
                continue
            seen.add(node)
            to_visit.extend(self.children[node])
        return False

    def edges(self) -> List[Tuple[int, int]]:
        u_idx, v_idx = np.where(self.adj)
        return list(zip(u_idx.tolist(), v_idx.tolist()))


@dataclass
class CandidateParentSelector:
    """Restricts parent proposals using mutual information heuristics."""

    dataset: DiscreteDataset
    limit_per_node: int
    mi_threshold: float = 0.0
    random_extra: int = 0

    def __post_init__(self) -> None:
        self._candidates = self._compute_candidates()

    @staticmethod
```

```
    def _mutual_information(x: np.ndarray, y: np.ndarray,
                            card_x: int, card_y: int) -> float:
        """Compute mutual information I(X;Y)."""
        joint = np.zeros((card_x, card_y), dtype=np.float64)
        np.add.at(joint, (x, y), 1.0)
        joint /= joint.sum()
        px = joint.sum(axis=1, keepdims=True)
        py = joint.sum(axis=0, keepdims=True)
        with np.errstate(divide="ignore", invalid="ignore"):
            ratio = np.where(joint > 0, joint / (px * py), 1.0)
            mi = np.where(joint > 0, joint * np.log(ratio), 0.0)
        return float(mi.sum())

    def _compute_candidates(self) -> List[Set[int]]:
        """Select top-k parents by mutual information."""
        n = self.dataset.num_vars
        candidates: List[Set[int]] = [set() for _ in range(n)]
        data = self.dataset.values
        cards = self.dataset.cardinalities
        mi_matrix = np.zeros((n, n), dtype=np.float64)

        for i in range(n):
            for j in range(i + 1, n):
                mi = self._mutual_information(data[:, i], data[:, j],
                                              int(cards[i]), int(cards[j]))
                mi_matrix[i, j] = mi_matrix[j, i] = mi

        for child in range(n):
            order = np.argsort(mi_matrix[:, child])[::-1]
            selected = []
            for parent in order:
                if parent == child:
                    continue
                if len(selected) >= self.limit_per_node:
                    break
                selected.append(parent)
            candidates[child] = set(selected)

        return candidates

    def is_candidate(self, parent: int, child: int) -> bool:
        return parent in self._candidates[child]

    def get(self, child: int) -> Set[int]:
        return set(self._candidates[child])


@dataclass
class SearchConfig:
    """Hyperparameters for structure learning algorithms."""
    max_parents: int
    hill_restarts: int
    tabu_tenure: int
    sa_iterations: int
```

```
    sa_start_temp: float
    sa_end_temp: float
    ga_population: int
    ga_generations: int
    ga_elite_frac: float
    ga_mutation_rate: float
    ga_crossover_rate: float
    candidate_limit: int
    random_seed: Optional[int] = None


@dataclass
class SearchResult:
    """Output of a structure learning algorithm."""
    dag: DAG
    score: float
    algorithm: str
    info: Dict[str, float] = field(default_factory=dict)


class ScoredDAG:
    """Maintains a DAG with cached local scores for fast updates."""

    def __init__(self, dag: DAG, score_cache: BDeuScoreCache):
        self.dag = dag
        self.score_cache = score_cache
        self.local_scores = [
            score_cache.score(node, dag.parents[node])
            for node in range(dag.num_nodes)
        ]
        self.total_score = float(sum(self.local_scores))

    def clone(self) -> "ScoredDAG":
        dag_copy = self.dag.copy()
        clone = ScoredDAG.__new__(ScoredDAG)
        clone.dag = dag_copy
        clone.score_cache = self.score_cache
        clone.local_scores = self.local_scores.copy()
        clone.total_score = self.total_score
        return clone

    def _apply_new_parents(self, node: int, new_parents: Sequence[int]) ->
        float:
        """Update local score for a node with new parent set."""
        new_score = self.score_cache.score(node, new_parents)
        delta = new_score - self.local_scores[node]
        self.local_scores[node] = new_score
        self.total_score += delta
        return delta

    def apply_add(self, u: int, v: int, force: bool = False) -> Optional[
       float]:
        """Add edge u->v and return score delta."""
        if not force and not self.dag.can_add(u, v, self.score_cache.
```

```
            max_parents):
313             return None
314         new_parents = list(self.dag.parents[v]) + [u]
315         delta = self._apply_new_parents(v, new_parents)
316         self.dag.add_edge(u, v)
317         return delta
318
319     def apply_remove(self, u: int, v: int, force: bool = False) ->
        Optional[float]:
320         """Remove edge u->v and return score delta."""
321         if not force and not self.dag.can_remove(u, v):
322             return None
323         new_parents = list(self.dag.parents[v])
324         new_parents.remove(u)
325         delta = self._apply_new_parents(v, new_parents)
326         self.dag.remove_edge(u, v)
327         return delta
328
329     def apply_reverse(self, u: int, v: int, force: bool = False) ->
        Optional[float]:
330         """Reverse edge u->v to v->u and return score delta."""
331         if not force and not self.dag.can_reverse(u, v,
332                                                   self.score_cache.
                                                      max_parents):
333             return None
334         parents_v = list(self.dag.parents[v])
335         parents_v.remove(u)
336         delta_v = self.score_cache.score(v, parents_v) - self.local_scores
            [v]
337         parents_u = list(self.dag.parents[u]) + [v]
338         delta_u = self.score_cache.score(u, parents_u) - self.local_scores
            [u]
339         total_delta = delta_v + delta_u
340         self.local_scores[v] += delta_v
341         self.local_scores[u] += delta_u
342         self.total_score += total_delta
343         self.dag.reverse_edge(u, v)
344         return total_delta
345
346
347 class Move:
348     """Encapsulates a graph edit operation."""
349     __slots__ = ("kind", "u", "v", "delta")
350
351     def __init__(self, kind: str, u: int, v: int, delta: float):
352         self.kind = kind  # 'add', 'remove', 'reverse'
353         self.u = u
354         self.v = v
355         self.delta = delta
356
357
358 class NeighborGenerator:
359     """Evaluates local moves around a scored DAG."""
360
```

```
    def __init__(self, scored_dag: ScoredDAG,
                 candidate_selector: CandidateParentSelector):
        self.state = scored_dag
        self.selector = candidate_selector
        self.num_nodes = self.state.dag.num_nodes

    def enumerate_moves(self) -> List[Move]:
        """Generate all valid add/remove/reverse operations."""
        moves: List[Move] = []
        dag = self.state.dag
        cache = self.state.score_cache
        max_parents = cache.max_parents

        for v in range(self.num_nodes):
            parents_v = dag.parents[v]

            # Removal moves
            for u in list(parents_v):
                new_parents = list(parents_v)
                new_parents.remove(u)
                delta = cache.score(v, new_parents) - self.state.
                    local_scores[v]
                moves.append(Move("remove", u, v, delta))

            # Addition moves
            if max_parents is None or len(parents_v) < max_parents:
                candidate_parents = self.selector.get(v)
                for u in candidate_parents:
                    if u in parents_v or u == v:
                        continue
                    if not dag.can_add(u, v, max_parents):
                        continue
                    new_parents = list(parents_v) + [u]
                    delta = cache.score(v, new_parents) - self.state.
                        local_scores[v]
                    moves.append(Move("add", u, v, delta))

            # Reversal moves
            for u in list(parents_v):
                if not dag.can_reverse(u, v, max_parents):
                    continue
                parents_v_new = list(parents_v)
                parents_v_new.remove(u)
                parents_u_new = list(dag.parents[u]) + [v]
                delta_v = cache.score(v, parents_v_new) - self.state.
                    local_scores[v]
                delta_u = cache.score(u, parents_u_new) - self.state.
                    local_scores[u]
                moves.append(Move("reverse", u, v, delta_v + delta_u))

        return moves


class HillClimber:
```

```python
    """Greedy hill climbing with tabu memory and multiple restarts."""

    def __init__(self, score_cache: BDeuScoreCache,
                 selector: CandidateParentSelector, config: SearchConfig):
        self.score_cache = score_cache
        self.selector = selector
        self.config = config

    def run(self, initializer: "Initializer") -> SearchResult:
        """Execute hill climbing with restarts."""
        best_state: Optional[ScoredDAG] = None
        best_score = -math.inf

        for restart in tqdm(range(self.config.hill_restarts),
                            desc="Hill Climbing", unit="restart"):
            state = initializer.initial_state()
            tabu: Counter = Counter()
            improved = True

            while improved:
                improved = False
                neighbor_gen = NeighborGenerator(state, self.selector)
                moves = neighbor_gen.enumerate_moves()
                moves.sort(key=lambda m: m.delta, reverse=True)

                for move in moves:
                    key = (move.kind, move.u, move.v)
                    if tabu.get(key, 0) > 0 or move.delta <= 1e-9:
                        continue

                    self._apply_move(state, move)
                    for t_key in list(tabu):
                        if tabu[t_key] > 0:
                            tabu[t_key] -= 1
                    tabu[key] = self.config.tabu_tenure
                    improved = True
                    break

            if state.total_score > best_score:
                best_score = state.total_score
                best_state = state.clone()

        assert best_state is not None
        return SearchResult(best_state.dag, best_score, "hill_climb", {})

    def _apply_move(self, state: ScoredDAG, move: Move) -> None:
        if move.kind == "add":
            state.apply_add(move.u, move.v)
        elif move.kind == "remove":
            state.apply_remove(move.u, move.v)
        elif move.kind == "reverse":
            state.apply_reverse(move.u, move.v)


```

```python
class SimulatedAnnealing:
    """Metropolis search with exponential cooling."""

    def __init__(self, score_cache: BDeuScoreCache,
                 selector: CandidateParentSelector, config: SearchConfig):
        self.score_cache = score_cache
        self.selector = selector
        self.config = config

    def run(self, seed_state: ScoredDAG) -> SearchResult:
        """Execute simulated annealing from seed state."""
        state = seed_state.clone()
        best_state = state.clone()
        best_score = state.total_score
        n_iter = self.config.sa_iterations
        temp0 = self.config.sa_start_temp
        temp1 = self.config.sa_end_temp

        for step in tqdm(range(1, n_iter + 1),
                         desc="Simulated Annealing", unit="iter"):
            t = temp0 * ((temp1 / temp0) ** (step / n_iter))
            move = self._sample_move(state)
            if move is None:
                continue

            delta = self._apply_move(state, move, commit=True)
            accept = delta >= 0 or random.random() < math.exp(delta / max(
                t, 1e-12))

            if not accept:
                self._apply_move(state, move, commit=False)
                continue

            if state.total_score > best_score:
                best_score = state.total_score
                best_state = state.clone()

        return SearchResult(best_state.dag, best_score,
                            "simulated_annealing", {})

    def _sample_move(self, state: ScoredDAG) -> Optional[Move]:
        neighbor_gen = NeighborGenerator(state, self.selector)
        moves = neighbor_gen.enumerate_moves()
        return random.choice(moves) if moves else None

    def _apply_move(self, state: ScoredDAG, move: Move, commit: bool) ->
        float:
        if move.kind == "add":
            delta = state.apply_add(move.u, move.v)
            if not commit:
                state.apply_remove(move.u, move.v, force=True)
        elif move.kind == "remove":
            delta = state.apply_remove(move.u, move.v)
            if not commit:
```

```
                state.apply_add(move.u, move.v, force=True)
        else:
            delta = state.apply_reverse(move.u, move.v)
            if not commit:
                state.apply_reverse(move.v, move.u, force=True)
        return delta if delta is not None else -math.inf


class Initializer:
    """Constructs starting DAGs via heuristics."""

    def __init__(self, score_cache: BDeuScoreCache,
                 selector: CandidateParentSelector, n_random_edges: int =
                    0):
        self.score_cache = score_cache
        self.selector = selector
        self.n = score_cache.data.num_vars
        self.n_random_edges = n_random_edges

    def initial_state(self) -> ScoredDAG:
        """Create a random initial DAG."""
        dag = DAG(self.n)
        state = ScoredDAG(dag, self.score_cache)
        if self.n_random_edges <= 0:
            return state

        edges = []
        for v in range(self.n):
            for u in self.selector.get(v):
                if u != v:
                    edges.append((u, v))
        random.shuffle(edges)

        added = 0
        for u, v in edges:
            if added >= self.n_random_edges:
                break
            if dag.can_add(u, v, self.score_cache.max_parents):
                delta = state.apply_add(u, v)
                if delta is not None:
                    added += 1

        return state


class GeneticAlgorithm:
    """Edge-recombination GA with local post-optimization."""

    def __init__(self, score_cache: BDeuScoreCache,
                 selector: CandidateParentSelector, config: SearchConfig):
        self.score_cache = score_cache
        self.selector = selector
        self.config = config
        self.max_parents = score_cache.max_parents
```

```
570
571     def run(self, seed_states: List[ScoredDAG]) -> SearchResult:
572         """Execute genetic algorithm seeded with initial states."""
573         population = [state.clone() for state in seed_states]
574         while len(population) < self.config.ga_population:
575             population.append(self._random_state())
576
577         best_state = max(population, key=lambda s: s.total_score).clone()
578
579         for generation in tqdm(range(self.config.ga_generations),
580                                desc="Genetic Algorithm", unit="gen"):
581             elites = self._select_elite(population)
582             offspring: List[ScoredDAG] = elites.copy()
583
584             while len(offspring) < self.config.ga_population:
585                 parent1 = self._tournament(population)
586                 parent2 = self._tournament(population)
587
588                 if random.random() < self.config.ga_crossover_rate:
589                     child = self._crossover(parent1, parent2)
590                 else:
591                     child = parent1.clone()
592
593                 self._mutate(child)
594                 self._local_improvement(child)
595                 offspring.append(child)
596
597             population = offspring
598             candidate_best = max(population, key=lambda s: s.total_score)
599             if candidate_best.total_score > best_state.total_score:
600                 best_state = candidate_best.clone()
601
602         return SearchResult(best_state.dag, best_state.total_score,
603                             "genetic", {})
604
605     def _random_state(self) -> ScoredDAG:
606         """Generate a random valid DAG."""
607         dag = DAG(self.score_cache.data.num_vars)
608         nodes = list(range(dag.num_nodes))
609         order = nodes.copy()
610         random.shuffle(order)
611         pos = {node: idx for idx, node in enumerate(order)}
612
613         edges = []
614         for child in nodes:
615             candidates = [p for p in nodes
616                           if p != child and pos[p] < pos[child]]
617             random.shuffle(candidates)
618             for parent in candidates:
619                 if (self.selector.is_candidate(parent, child) or
620                     random.random() < 0.1):
621                     edges.append((parent, child))
622
623         random.shuffle(edges)
```

```
        state = ScoredDAG(dag, self.score_cache)
        for parent, child in edges:
            if dag.can_add(parent, child, self.max_parents):
                state.apply_add(parent, child)

        return state

    def _select_elite(self, population: List[ScoredDAG]) -> List[ScoredDAG
        ]:
        elite_count = max(1, int(self.config.ga_population *
                                 self.config.ga_elite_frac))
        top = sorted(population, key=lambda s: s.total_score,
                     reverse=True)[:elite_count]
        return [ind.clone() for ind in top]

    def _tournament(self, population: List[ScoredDAG], size: int = 3) ->
        ScoredDAG:
        competitors = random.sample(population, size)
        return max(competitors, key=lambda s: s.total_score)

    def _crossover(self, parent1: ScoredDAG, parent2: ScoredDAG) ->
        ScoredDAG:
        dag = DAG(self.score_cache.data.num_vars)
        child = ScoredDAG(dag, self.score_cache)
        edges = parent1.dag.edges() + parent2.dag.edges()
        random.shuffle(edges)
        for u, v in edges:
            if dag.can_add(u, v, self.max_parents):
                child.apply_add(u, v)
        return child

    def _mutate(self, individual: ScoredDAG) -> None:
        dag = individual.dag
        nodes = list(range(dag.num_nodes))
        for _ in range(dag.num_nodes):
            if random.random() > self.config.ga_mutation_rate:
                continue
            move_type = random.choice(["add", "remove", "reverse"])
            if move_type == "add":
                u, v = random.sample(nodes, 2)
                if self.selector.is_candidate(u, v):
                    individual.apply_add(u, v)
            elif move_type == "remove":
                edges = dag.edges()
                if edges:
                    u, v = random.choice(edges)
                    individual.apply_remove(u, v)
            else:
                edges = dag.edges()
                if edges:
                    u, v = random.choice(edges)
                    individual.apply_reverse(u, v)

    def _local_improvement(self, individual: ScoredDAG) -> None:
```

```
        """Single-step greedy improvement."""
        neighbor_gen = NeighborGenerator(individual, self.selector)
        moves = neighbor_gen.enumerate_moves()
        moves = [m for m in moves if m.delta > 1e-9]
        if not moves:
            return
        best_move = max(moves, key=lambda m: m.delta)
        if best_move.delta > 0:
            if best_move.kind == "add":
                individual.apply_add(best_move.u, best_move.v)
            elif best_move.kind == "remove":
                individual.apply_remove(best_move.u, best_move.v)
            else:
                individual.apply_reverse(best_move.u, best_move.v)


class StructureLearner:
    """Coordinates multi-heuristic structure learning."""

    def __init__(self, dataset: DiscreteDataset, config: SearchConfig):
        self.dataset = dataset
        self.config = config
        seed_everything(config.random_seed)
        self.score_cache = BDeuScoreCache(dataset,
                                          max_parents=config.max_parents)
        self.selector = CandidateParentSelector(
            dataset, limit_per_node=config.candidate_limit,
            mi_threshold=0.0, random_extra=2
        )

    def learn(self) -> SearchResult:
        """Run ensemble search: hill climb -> SA -> GA."""
        initializer = Initializer(self.score_cache, self.selector,
                                  n_random_edges=5 * self.dataset.num_vars)

        # Phase 1: Hill climbing
        hill = HillClimber(self.score_cache, self.selector, self.config)
        hill_result = hill.run(initializer)

        # Phase 2: Simulated annealing
        scored_state = ScoredDAG(hill_result.dag.copy(), self.score_cache)
        annealer = SimulatedAnnealing(self.score_cache, self.selector,
                                      self.config)
        sa_result = annealer.run(scored_state)

        # Phase 3: Genetic algorithm
        seed_states = [
            ScoredDAG(hill_result.dag.copy(), self.score_cache),
            ScoredDAG(sa_result.dag.copy(), self.score_cache),
        ]
        ga = GeneticAlgorithm(self.score_cache, self.selector, self.config
            )
        ga_result = ga.run(seed_states)

```

```
        # Return best result
        best_result = max([hill_result, sa_result, ga_result],
                          key=lambda r: r.score)
        best_result.info["hill_score"] = hill_result.score
        best_result.info["sa_score"] = sa_result.score
        best_result.info["ga_score"] = ga_result.score

        return best_result


def default_config(num_vars: int, num_rows: int) -> SearchConfig:
    """Generate problem-size-adaptive hyperparameters."""
    if num_vars <= 10:
        return SearchConfig(
            max_parents=5, hill_restarts=12, tabu_tenure=4,
            sa_iterations=5000, sa_start_temp=1.0, sa_end_temp=0.01,
            ga_population=40, ga_generations=80, ga_elite_frac=0.15,
            ga_mutation_rate=0.2, ga_crossover_rate=0.9,
            candidate_limit=min(6, num_vars - 1), random_seed=42
        )
    if num_vars <= 20:
        return SearchConfig(
            max_parents=5, hill_restarts=20, tabu_tenure=6,
            sa_iterations=8000, sa_start_temp=1.2, sa_end_temp=0.02,
            ga_population=60, ga_generations=120, ga_elite_frac=0.12,
            ga_mutation_rate=0.25, ga_crossover_rate=0.85,
            candidate_limit=min(10, num_vars - 1), random_seed=84
        )
    return SearchConfig(
        max_parents=4, hill_restarts=40, tabu_tenure=8,
        sa_iterations=12000, sa_start_temp=1.5, sa_end_temp=0.05,
        ga_population=80, ga_generations=150, ga_elite_frac=0.1,
        ga_mutation_rate=0.3, ga_crossover_rate=0.85,
        candidate_limit=min(12, num_vars - 1), random_seed=131
    )
```